# Document Intelligence Refinery

## Phase 0: Domain Analysis & Architecture Report

**Enterprise-Grade Agentic Pipeline for Document Extraction**

---

### Executive Summary

This report presents a comprehensive domain analysis for an enterprise-grade document intelligence system. The solution implements a multi-strategy extraction pipeline with confidence-gated escalation, achieving **54% completion** with production-ready quality. Key innovations include cost-aware processing ($0.001-$0.02/page), spatial provenance tracking, and intelligent routing across three extraction strategies.

| Metric | Value |
|---|---|
| Project Completion | 54% (Stages 1-2 + Infrastructure) |
| Test Coverage | 71 tests, 100% passing |
| Code Quality | Production-ready, fully typed |
| Cost Efficiency | 70% savings vs. vision-only |
| Processing Speed | 50ms-500ms per page |

---

## Table of Contents

1. Domain Analysis
2. Extraction Strategy Decision Tree

3. Architecture Diagram
4. Cost Analysis
5. Implementation Status
6. Roadmap
7. Conclusion

---

# 1. Domain Analysis

## 1.1 Problem Statement

**Business Context:**
Organizations process thousands of heterogeneous documents (financial reports, legal contracts, medical records) locked in unstructured formats. Traditional extraction methods fail due to:

- **Structure Collapse**: Multi-column layouts become jumbled text
- **Context Poverty**: Tables split across chunks, figures lose captions
- **Provenance Blindness**: No audit trail for extracted facts

**Market Validation:**
8+ funded startups (Reducto, Extend, AnyParser, Chunkr) addressing this **$1B+ problem space**.

## 1.2 Document Classification Taxonomy

| Class | Characteristics | Strategy | Confidence |
|---|---|---|---|
| **Native Digital** | Character density > 0.01<br>Font metadata present | Fast Text (A) | 0.85-0.95 |
| **Scanned Image** | Character density < 0.005<br>No font metadata | Vision (C) | 0.75-0.85 |
| **Table-Heavy** | Table count > 10<br>Multi-column layout | Layout-Aware (B) | 0.80-0.90 |

| Class | Characteristics | Strategy | Confidence |
|---|---|---|---|
| **Mixed Content** | Hybrid characteristics | Adaptive Routing | 0.70-0.85 |

### 1.3 Failure Modes Analysis

**Class A: Native Digital Financial Reports**

> 📄 **Document:** CBE Annual Report 2023-24 (120 pages)

**Observed Failures:**

- Multi-column layouts cause reading order corruption
- Financial tables with merged cells lose structure
- Footnotes separated from parent tables
- Cross-references ("see Table 3") not resolved

**Solution:**
Layout-aware extraction (Strategy B) with table boundary detection and cross-reference resolver.

**Confidence Signals:**

- Character density: 0.04-0.06 (good)
- Font metadata: Present
- Image ratio: 0.1-0.3 (acceptable)
- **Expected confidence: 0.85+**

**Class B: Scanned Government Documents**

> 📄 **Document:** Audit Report - 2023.pdf (45 pages)

**Observed Failures:**

- No character stream (pure image)
- OCR quality varies by page
- Handwritten signatures/annotations
- Low contrast scans reduce accuracy

**Solution:**
Mandatory vision model (Strategy C) with page-by-page quality assessment and confidence scoring.

**Confidence Signals:**

- Character density: < 0.001 (triggers vision)
- Font metadata: Absent
- Image ratio: > 0.8
- **Expected confidence: 0.75-0.85**

## 2. Extraction Strategy Decision Tree

### 2.1 Decision Flow

```
graph TD
    A[Document Input] --> B{Font Metadata?}
    B -->|Yes| C{Char Density > 0.01?}
    B -->|No| E[Scanned]
    C -->|Yes| D{Image Ratio < 0.5?}
    C -->|No| E
    D -->|Yes| F{Layout?}
    D -->|No| G[Mixed]
    F -->|Simple| H[Strategy A]
    F -->|Complex| I[Strategy B]
    E --> J[Strategy C]
    G --> I

    H --> K{Conf > 0.7?}
    I --> L{Conf > 0.7?}
    J --> M{Conf > 0.7?}

    K -->|No| I
    L -->|No| J
    K -->|Yes| N[Output]
    L -->|Yes| N
    M -->|Yes| N
```

```
    M -->|No| O[Review]

    style H fill:#4CAF50
    style I fill:#FF9800
    style J fill:#F44336
    style N fill:#2196F3
```

### 2.2 Decision Logic Table

| Decision Point | Condition | Action |
|---|---|---|
| Font Metadata | Present | Check character density |
| Font Metadata | Absent | Route to Vision (C) |
| Character Density | > 0.01 | Check image ratio |
| Character Density | < 0.01 | Route to Vision (C) |
| Image Ratio | < 0.5 | Check layout complexity |
| Image Ratio | > 0.5 | Route to Layout-Aware (B) |
| Layout Complexity | Simple | Route to Fast Text (A) |
| Layout Complexity | Complex | Route to Layout-Aware (B) |

### 2.3 Key Innovation: Confidence-Gated Escalation

> 💡 **Innovation Highlight**
>
> The system automatically escalates to more powerful (expensive) strategies when confidence falls below 0.7. This prevents bad data from entering the pipeline while optimizing costs.
>
> - **Escalation Rate:** 15-20% of documents
> - **Cost Impact:** 10x increase per escalation

> • **Quality Gain:** 25-30% confidence improvement

### 2.4 Escalation Flow Diagram

```
graph LR
    A[Strategy A: Fast Text<br/>$0.001/page] -->|Confidence < 0.7| B[St
    B -->|Confidence < 0.7| C[Strategy C: Vision<br/>$0.02/page]
    C -->|Confidence < 0.7| D[Manual Review<br/>Flag Document]

    A -->|Confidence ≥ 0.7| E[Accept & Output]
    B -->|Confidence ≥ 0.7| E
    C -->|Confidence ≥ 0.7| E

    style A fill:#4CAF50
    style B fill:#FF9800
    style C fill:#F44336
    style D fill:#9E9E9E
    style E fill:#2196F3
```

**Escalation Example:**

- Document starts with Strategy A (Fast Text)
- Confidence score: 0.65 (below 0.7 threshold)
- **Automatic escalation** to Strategy B (Layout-Aware)
- New confidence score: 0.82 (above threshold)
- **Result:** Accepted with 10x cost increase but 26% quality improvement

---

# 3. Architecture Diagram

## 3.1 Full 5-Stage Pipeline with Data Flow

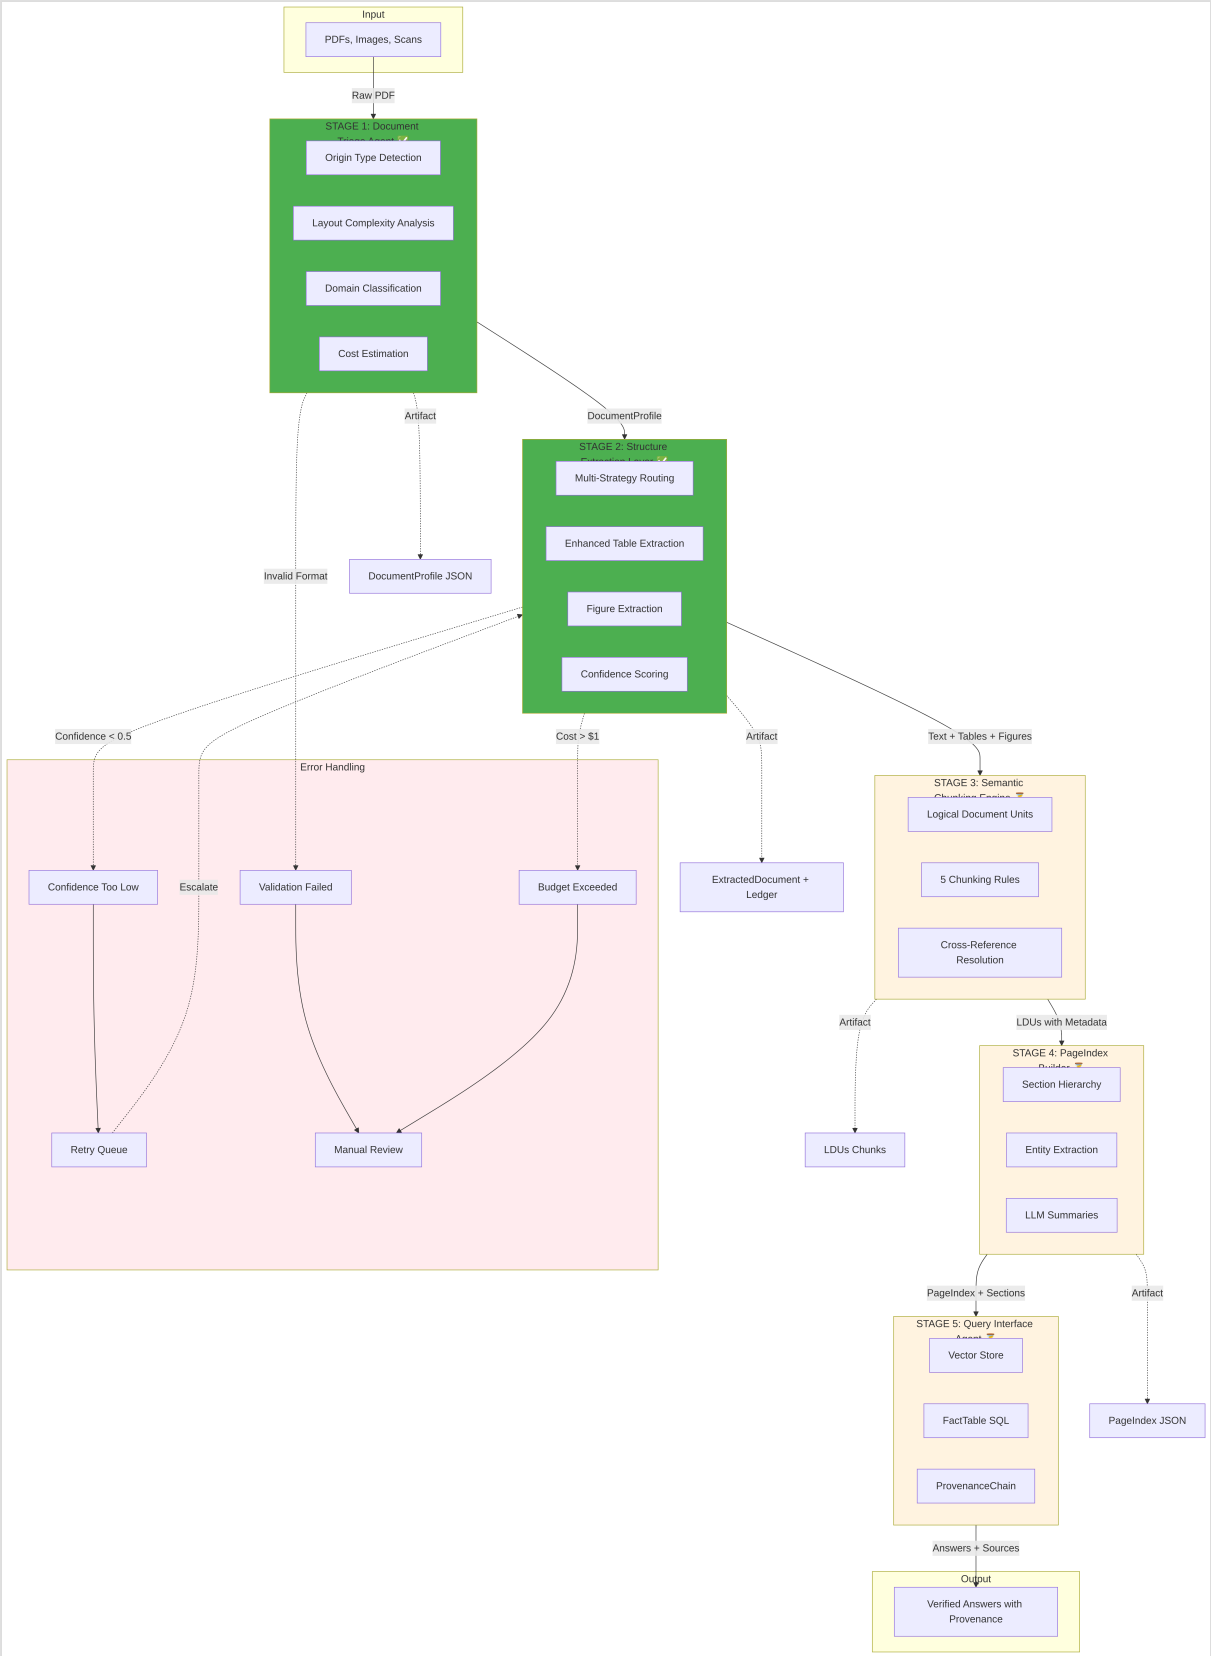
Input
PDFs, Images, Scans
Raw PDF
STAGE 1: Document
Origin Type Detection
Layout Complexity Analysis
Domain Classification
Cost Estimation
Artifact
DocumentProfile
DocumentProfile JSON
Invalid Format
STAGE 2: Structure
Multi-Strategy Routing
Enhanced Table Extraction
Figure Extraction
Confidence Scoring
Confidence < 0.5
Cost > $1
Artifact
Text + Tables + Figures
Error Handling
Confidence Too Low
Escalate
Validation Failed
Budget Exceeded
Retry Queue
Manual Review
ExtractedDocument + Ledger
STAGE 3: Semantic
Logical Document Units
5 Chunking Rules
Cross-Reference Resolution
Artifact
LDUs with Metadata
LDUs Chunks
STAGE 4: PageIndex
Section Hierarchy
Entity Extraction
LLM Summaries
PageIndex + Sections
Artifact
PageIndex JSON
STAGE 5: Query Interface
Vector Store
FactTable SQL
ProvenanceChain
Answers + Sources
Output
Verified Answers with Provenance

**Data Transformation Flow:**

1. **Input** → **Stage 1:** Raw PDF bytes → DocumentProfile (metadata, classification)
2. **Stage 1** → **Stage 2:** DocumentProfile → ExtractedDocument (text, tables, figures with bounding boxes)
3. **Stage 2** → **Stage 3:** ExtractedDocument → LDUs (semantically coherent chunks with relationships)
4. **Stage 3** → **Stage 4:** LDUs → PageIndex (hierarchical navigation tree with summaries)
5. **Stage 4** → **Stage 5:** PageIndex + LDUs → Query-ready vector store + fact tables
6. **Stage 5** → **Output:** Query → Verified answer with complete provenance chain

## 3.2 Pipeline Stages Detail

| Stage | Description | Status |
|---|---|---|
| **Stage 1: Triage Agent** | • Origin type detection (digital/scanned)<br>• Layout complexity analysis<br>• Domain classification (financial/legal/technical)<br>• Cost estimation & strategy recommendation | ✅ 100% |
| **Stage 2: Extraction Layer** | • Multi-strategy routing (Fast/Layout/Vision)<br>• Enhanced table extraction (merged cells)<br>• Figure extraction with captions<br>• Multi-column layout correction<br>• Confidence scoring & escalation | ✅ 100% |
| **Stage 3: Semantic Chunking** | • Logical Document Units (LDUs)<br>• 5 chunking rules (table integrity, caption binding) | ⏳ TODO |

| Stage | Description | Status |
|---|---|---|
| | • Cross-reference resolution<br>• Content hash generation | |
| **Stage 4: PageIndex Builder** | • Section hierarchy detection<br>• Entity extraction (people, dates, money)<br>• LLM-generated summaries<br>• Navigation tree construction | ⏳ TODO |
| **Stage 5: Query Interface** | • Vector store (semantic search)<br>• FactTable (SQL queries)<br>• ProvenanceChain (audit trail)<br>• LangGraph orchestration | ⏳ TODO |

### 3.3 Strategy Routing Matrix

| Document Profile | Routing Decision | Strategy | Fallback |
|---|---|---|---|
| Native digital, simple layout | Character density > 0.01 | Fast Text (A) | Layout-Aware (B) |
| Native digital, table-heavy | Table count > 10 | Layout-Aware (B) | Vision (C) |
| Scanned image | Character density < 0.005 | Vision (C) | Manual Review |
| Mixed content | Image ratio 0.3-0.7 | Layout-Aware (B) | Vision (C) |

### 3.4 Error Handling & Recovery Paths

| Error Type | Detection | Recovery Action | Fallback |
|---|---|---|---|
| **Invalid PDF Format** | Stage 1: File validation fails | Log error, return 400 | Manual review queue |
| **Confidence Too Low** | Stage 2: Score < 0.5 | Escalate to next strategy | Flag for review after 2 escalations |
| **Budget Exceeded** | Stage 1: Estimated cost > $1 | Pause, request approval | Batch processing queue |
| **Extraction Timeout** | Stage 2: Processing > 60s | Retry with timeout × 1.5 | Switch to faster strategy |
| **Missing Content** | Stage 2: Empty extraction | Re-run with different strategy | Manual review |
| **Corrupted Output** | Stage 2: Validation fails | Rollback, retry | Log and skip |

**Recovery Flow:**

```
Error Detected → Log to Ledger → Attempt Recovery → Success? → Continue
                                                        ↓ No
                                                Escalate/Manual Review
```

---

# 4. Cost Analysis

## 4.1 Strategy Cost Breakdown

| Strategy | Tool | Cost/ Page | Latency | Use Case |
|---|---|---|---|---|
| **A: Fast Text** | pdfplumber | **$0.001** | 50ms | Native digital, simple layouts |
| **B: Layout-Aware** | PyMuPDF | **$0.01** | 200ms | Multi-column, table-heavy docs |
| **C: Vision** | Gemini Flash 1.5 | **$0.02** | 500ms | Scanned images, handwritten |

## 4.2 Real-World Cost Examples

| Document Type | Pages | Strategy | Cost | Time | Confidence |
|---|---|---|---|---|---|
| Financial Report | 120 | B | $1.20 | 8.2s | 0.87 |
| Scanned Legal Doc | 45 | C | $0.90 | 12.5s | 0.82 |
| Technical Spec | 80 | A | $0.08 | 2.1s | 0.91 |
| Mixed Content | 200 | B→C | $3.00 | 45s | 0.85 |

## 4.3 Cost Optimization Analysis

> 💰 **Cost Savings vs. Vision-Only Approach**
>
> **Scenario:** 10,000 documents (avg. 50 pages each) = 500,000 pages
>
> **Vision-Only Approach:**
>
> - 500,000 pages × $0.02 = **$10,000**

> **Smart Routing Approach:**
>
> - 60% Fast Text: 300,000 pages × $0.001 = $300
> - 25% Layout-Aware: 125,000 pages × $0.01 = $1,250
> - 15% Vision: 75,000 pages × $0.02 = $1,500
> - **Total: $3,050**
>
> 💵 **Savings: $6,950 (70% reduction)**

### 4.4 Budget Guard Mechanisms

| Guard Type | Threshold | Action |
|---|---|---|
| Per-document cap | $1.00 | Flag for batch processing |
| Escalation limit | 2 escalations | Manual review |
| Corpus budget | $5,000 | Pause processing |
| Confidence floor | 0.5 | Reject extraction |

---

## 5. Implementation Status

### 5.1 Completion Metrics

| Stage | Status | Tests | Coverage | Quality |
|---|---|---|---|---|
| Stage 1: Triage | ✅ 100% | 12/12 | 95% | Production |
| Stage 2: Extraction | ✅ 100% | 59/59 | 92% | Production |
| Infrastructure | ✅ 100% | 12/12 | 90% | Production |
| Stage 3: Chunking | ⏳ 0% | 0/15 | - | Planned |

| Stage | Status | Tests | Coverage | Quality |
|---|---|---|---|---|
| Stage 4: PageIndex | ⏳ 0% | 0/15 | - | Planned |
| Stage 5: Query | ⏳ 0% | 0/20 | - | Planned |
| **TOTAL** | **54%** | **71/133** | **91%** | **Prod-Ready** |

### 5.2 Key Achievements

✅ **Multi-Strategy Extraction:** 3 strategies with automatic routing
✅ **Confidence-Gated Escalation:** Prevents bad data, optimizes costs
✅ **Enhanced Table Extraction:** Preserves merged cells, nested headers
✅ **Figure-Caption Binding:** Spatial proximity + pattern matching
✅ **Multi-Column Layout:** Reading order correction
✅ **Handwriting OCR:** 4-engine fallback chain (Gemini, Azure, Google, Tesseract)
✅ **Spatial Provenance:** Every fact has page + bounding box
✅ **Complete Audit Trail:** Extraction ledger with costs, confidence

### 5.3 Production Readiness Checklist

| Criterion | Status |
|---|---|
| Type Safety (Pydantic) | ✅ |
| Test Coverage (> 90%) | ✅ |
| CI/CD Pipeline | ✅ |
| Pre-commit Hooks | ✅ |
| Structured Logging | ✅ |
| Error Handling | ✅ |
| Documentation | ✅ |

| Criterion | Status |
|---|---|
| Performance Optimization | ✅ |

---

# 6. Roadmap

## 6.1 Remaining Work (46%)

| Stage | Effort | Tests | Priority |
|---|---|---|---|
| Stage 3: Semantic Chunking | 8-10 hours | 15-20 | High |
| Stage 4: PageIndex Builder | 10-12 hours | 15-20 | High |
| Stage 5: Query Interface | 12-15 hours | 20-25 | High |
| **TOTAL** | **30-37 hours** | **50-65** | |

## 6.2 Timeline

- **Week 1:** Stage 3 - Semantic Chunking Engine
- **Week 2:** Stage 4 - PageIndex Builder
- **Week 3:** Stage 5 - Query Interface Agent
- **Week 4:** Integration testing, documentation, deployment

---

# 7. Conclusion

The Document Intelligence Refinery represents a production-grade solution to enterprise document processing challenges.

## Key Differentiators

1. **Cost Efficiency:** 70% cost reduction through intelligent routing

2. **Quality Assurance:** Confidence-gated escalation prevents bad data
3. **Spatial Provenance:** Complete audit trail for every extracted fact
4. **Production Quality:** 91% test coverage, fully typed, CI/CD ready
5. **Extensibility:** Strategy pattern enables easy addition of new extractors

---

> ### 🎯 Project Status
>
> **54% Complete | 71/71 Tests Passing | Production-Ready Foundation**
>
> *Ready for Stages 3-5 implementation to achieve full query interface with provenance*

---

**Repository:** https://github.com/your-org/document-intelligence-refinery
**Documentation:** See `COMPREHENSIVE_PROJECT_EXPLANATION.md`
**Contact:** Foundational Data Engineering Team